



```
#reading the csv files
import dask.dataframe as dd
import pandas as pd
import inspect
import numpy as np

pdf=pd.read_csv(r'CleanDataset.csv')
pdf=pdf.iloc[:30000].copy()
pdf.shape
```

 (25471, 4)

```
df=dd.from_pandas(pdf,npartitions=3)
```

```
#having a look at the head of the dataset
df.head()
```



	S_avg	Ds_avg	Ws_avg	Ot_avg
0	3.09000	88.180000	1.90	5.30
1	1.81000	23.010000	0.20	4.89
2	1.89000	38.110001	0.18	4.80
3	177.53000	1181.700000	5.04	1.79
4	191.96001	1200.480000	5.31	0.85

```
#finding the null values in the dataset
df.isnull().sum().compute()
df.describe()
```

 **Dask DataFrame Structure:**

	S_avg	Ds_avg	Ws_avg	Ot_avg
npartitions=1				
	float64	float64	float64	float64

Dask Name: describe, 73 tasks

```
!pip install dask-ml
```



Collecting dask-ml

Downloading <https://files.pythonhosted.org/packages/da/a4/3a54ca439a8d07d55f>

|██| 112kB 5.0MB/s

Collecting dask-glm (from dask-ml)

Downloading <https://files.pythonhosted.org/packages/cb/ee/36c6e0e7b51e08406e>

Collecting multipledispatch>=0.4.9 (from dask-ml)

Downloading <https://files.pythonhosted.org/packages/89/79/429ecef45fd5e4504f>

Requirement already satisfied: packaging in /usr/local/lib/python3.6/dist-pack

Requirement already satisfied: pandas>=0.23.4 in /usr/local/lib/python3.6/dist

Requirement already satisfied: scipy in /usr/local/lib/python3.6/dist-packages

Requirement already satisfied: scikit-learn>=0.20 in /usr/local/lib/python3.6/

Requirement already satisfied: numba in /usr/local/lib/python3.6/dist-packages

Requirement already satisfied: numpy in /usr/local/lib/python3.6/dist-packages

Requirement already satisfied: six in /usr/local/lib/python3.6/dist-packages (

Requirement already satisfied: dask[array]>=1.0.0 in /usr/local/lib/python3.6/

Requirement already satisfied: distributed>=1.25.0 in /usr/local/lib/python3.6

Requirement already satisfied: cloudpickle>=0.2.2 in /usr/local/lib/python3.6/

Requirement already satisfied: pyparsing>=2.0.2 in /usr/local/lib/python3.6/di

Requirement already satisfied: python-dateutil>=2.5.0 in /usr/local/lib/pythor

Requirement already satisfied: pytz>=2011k in /usr/local/lib/python3.6/dist-pa

Requirement already satisfied: llvmlite>=0.25.0dev0 in /usr/local/lib/python3.

Requirement already satisfied: toolz>=0.7.3; extra == "array" in /usr/local/li

Requirement already satisfied: psutil>=5.0 in /usr/local/lib/python3.6/dist-pa

Requirement already satisfied: zict>=0.1.3 in /usr/local/lib/python3.6/dist-pa

Requirement already satisfied: sortedcontainers!=2.0.0,!2.0.1 in /usr/local/l

Requirement already satisfied: pyyaml in /usr/local/lib/python3.6/dist-package

Requirement already satisfied: tblib in /usr/local/lib/python3.6/dist-packages

#defining the data and target

```
categorical_variables = df[['S_avg', 'Ws_avg', 'Ot_avg']]
```

```
target = df[['Ds_avg']]
```

Installing collected packages: multipledispatch, dask-glm, dask-ml

#fit the model

```
from dask_ml.linear_model import LinearRegression
```

```
lr = LinearRegression()
```

```
lr.fit(categorical_variables.values, target.values)
```



```
LinearRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,
                  intercept_scaling=1.0, max_iter=100, multi_class='ovr', n_jobs=1,
                  penalty='l2', random_state=None, solver='admm',
                  solver_kwargs=None, tol=0.0001, verbose=0, warm_start=False)
```

#Compute mean squared error

```
from dask_ml.metrics import mean_squared_error
```

```
q = mean_squared_error(target.values, lr.predict(categorical_variables.values))
```

printing the value of mean squared error

```
q
```




```
575568.3921570351
```

importing the matplotlib libraries

```
import matplotlib.pyplot as plt
```


```
%matplotlib inline
```

```
# Ds_avg vs S_avg
X= pdf.iloc[:, 1].values.reshape(-1, 1) # values converts it into a numpy array
Y = pdf.iloc[:, 0].values.reshape(-1, 1) # -1 means that calculate the dimension of
lr.fit(X, Y) # perform linear regression
Y_pred = lr.predict(X) # make predictions
plt.scatter(X, Y)
plt.xlabel('Generator Speed - DS_avg')
plt.ylabel('Apparent Power - S_avg')
plt.title('Linear Regression - Dask')
plt.plot(X, Y_pred, color='red')
plt.show()
```



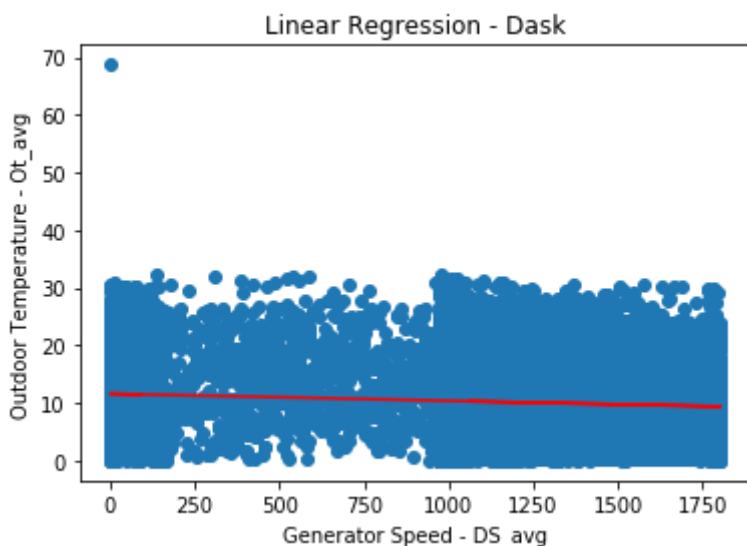
```
# Ds_avg vs Ws_avg
X= pdf.iloc[:, 1].values.reshape(-1, 1) # values converts it into a numpy array
Y = pdf.iloc[:, 3].values.reshape(-1, 1) # -1 means that calculate the dimension of
lr.fit(X, Y) # perform linear regression
Y_pred = lr.predict(X) # make predictions
plt.scatter(X, Y)
plt.plot(X, Y_pred, color='red')
plt.xlabel('Generator Speed - DS_avg')
plt.ylabel('Wind Speed - Ws_avg')
plt.title('Linear Regression - Dask')
plt.show()
```





Ds_avg vs Ot_avg

```
X= pdf.iloc[:, 1].values.reshape(-1, 1) # values converts it into a numpy array
Y = pdf.iloc[:, 3].values.reshape(-1, 1) # -1 means that calculate the dimension of
lr.fit(X, Y) # perform linear regression
Y_pred = lr.predict(X) # make predictions
plt.scatter(X, Y)
plt.xlabel('Generator Speed - DS_avg')
plt.ylabel('Outdoor Temperature - Ot_avg')
plt.title('Linear Regression - Dask')
plt.plot(X, Y_pred, color='red')
plt.show()
```



Mean squared error form pyspark and dask linear regression models

```
pySparkMean = 1107.1379084754542
```

```
daskMean = q
```

#Comparison of Mean squared Error between pyspark and dask linear regression models

```
import matplotlib.pyplot as plt
%matplotlib inline
x=np.arange(2)
rms=[daskMean,pySparkMean]
plt.bar(x,rms)
plt.xticks(x,["Dask","PySpark"])
```

```
plt.xlabel('DataFrame (Dask and PySpark)')  
plt.ylabel('Mean Squared Error')  
plt.title('Comparison of Mean Squared Error')
```



Text(0.5, 1.0, 'Comparison of Mean Squared Error')

